

Ollama on Windows — Home Inference Server

Offload medycabrain's AI workload from the VPS to your Windows laptop

Setup guide · medycabrain · July 2026

Windows laptop runs the model → **Tailscale** private link → **VPS (medycabrain)**

The laptop runs Ollama (the Qwen 7B model) as an always-on background service. The VPS sends it the AI jobs — cluster naming, content ideas, Q&A — over a private Tailscale link and gets the answers back. The VPS keeps doing everything else (scraping, transcription, the web app); only the heavy LLM work moves to the laptop.

The one hard rule: the laptop must stay **powered on and plugged in**. Closing the lid, turning off the screen, locking, or logging out are all fine with the settings below — but a full shutdown or power loss stops the server until it boots back up. Keep it plugged in, somewhere ventilated.

1 Stop the laptop from sleeping **MUST DO**

A sleeping laptop can't serve. Make it stay awake even with the lid closed (the screen can still turn off — that doesn't stop computation).

Point & click

- **Settings** → **System** → **Power & battery** → **Screen and sleep** → set “When plugged in, put my device to sleep after” = **Never**.
- **Control Panel** → **Hardware and Sound** → **Power Options** → **Choose what closing the lid does** → “When I close the lid (Plugged in) = **Do nothing**”.

Or, one admin terminal (faster)

Right-click Start → **Terminal (Admin)**, then:

```
:: never sleep / hibernate while plugged in
powercfg /change standby-timeout-ac 0
powercfg /change hibernate-timeout-ac 0
:: closing the lid does nothing (plugged in)
powercfg /setacvalueindex SCHEME_CURRENT SUB_BUTTONS LIDACTION 0
powercfg /setactive SCHEME_CURRENT
```

2 Install Ollama & open it to the network **MUST DO**

1. Download and install Ollama for Windows from ollama.com/download.
2. Set two **system environment variables** (Start → “Edit the system environment variables” → Environment Variables → New, under *System variables*):

Variable	Value	Why
OLLAMA_HOST	0.0.0.0	Lets the VPS reach it (not just this PC)

OLLAMA_KEEP_ALIVE	-1	Keeps the model loaded — no reload lag per call
-------------------	----	---

3. Sign out and back in (or reboot) so the variables take effect.

3 Download the model **MUST DO**

In a terminal, pull the same quantized model the VPS uses (~4.7 GB, one-time download):

```
ollama pull qwen2.5:7b-instruct-q4_K_M

:: quick test that it answers
ollama run qwen2.5:7b-instruct-q4_K_M "Rispondi in italiano: cos'è la menopausa?"
```

Tip: the laptop's GPU makes this much faster than the shared VPS CPU. Even a modest laptop GPU will beat the current setup comfortably.

4 Run Ollama as a service — survives logoff & reboot **RECOMMENDED**

The normal Ollama tray icon stops when you **log out**. To keep it serving no matter what happens to your Windows session, run it as a **service** with NSSM (free, tiny).

1. Download NSSM from nssm.cc/download, unzip it, and note the path to `nssm.exe` (use the `win64` one).
2. First, stop the auto-started tray Ollama (right-click its tray icon → Quit), so the port is free.
3. In an **admin terminal**, from the folder with `nssm.exe`:

```
:: adjust the path if Ollama installed elsewhere
nssm install Ollama "%LOCALAPPDATA%\Programs\Ollama\ollama.exe" serve
nssm set Ollama AppEnvironmentExtra OLLAMA_HOST=0.0.0.0 OLLAMA_KEEP_ALIVE=-1
nssm set Ollama Start SERVICE_AUTO_START
nssm start Ollama
```

Now Ollama starts at boot (before anyone logs in) and keeps running through lock, logoff, and Windows-Update reboots. To check or stop it later: `nssm status Ollama` · `nssm restart Ollama` · `nssm stop Ollama` .

5 Connect laptop and VPS with Tailscale **MUST DO**

Your laptop is behind your home router, so the VPS can't reach it directly. Tailscale builds a tiny private network between the two — no port-forwarding, works from anywhere.

1. Install Tailscale on the **laptop** (tailscale.com/download) and sign in.
2. Install Tailscale on the **VPS** (I can do this — one command) and sign in with the **same account**.
3. On the laptop, note its Tailscale address — a `100.x.y.z` number shown in the Tailscale menu. That's the stable address the VPS will call.

Tailscale also runs as a background service, so it too survives logoff and reboots.

6 Point the VPS at the laptop I DO THIS

This part is on me — a one-line change on the VPS, because the model address is already configurable. Once you give me the laptop's Tailscale IP, I set:

```
# on the VPS: /home/giuseppe/projects/medycabrain/BEC/.env
OLLAMA_URL=http://100.x.y.z:11434
```

and restart the backend. I'll also add a **health check**: if the laptop is ever unreachable (off, or off the network), the VPS automatically falls back to running the model locally, so nothing breaks — it just gets slower again until the laptop is back.

7 Verify it works

From the laptop, confirm the service is answering on the network:

```
:: should list the model
curl http://localhost:11434/api/tags
```

From the VPS side (I'll run this), a call to the laptop's Tailscale IP should return the model list and a fast generation. When it's working, the medycabrain "Generate ideas" and cluster-naming jobs finish in seconds instead of minutes.

8 Troubleshooting & rollback

Symptom	Fix
VPS can't reach the laptop	Check Tailscale is connected on both (green). Confirm <code>OLLAMA_HOST=0.0.0.0</code> is set and the service was restarted after setting it.
Works, then stops after you log out	You're running the tray app, not the NSSM service. Do step 4.
Laptop keeps sleeping	Re-check step 1; some laptops also sleep on "balanced" power plan — set the plan to keep the system awake when plugged in.
Model reloads slowly each call	Set <code>OLLAMA_KEEP_ALIVE=-1</code> (step 2) and restart the service.
Want to go back to VPS-only	Tell me — I revert one env line and the VPS runs the model itself again. No other changes.

End state: laptop plugged in, lid closed, screen off, logged out → still serving the VPS. A reboot recovers on its own. Only a full power-off stops it.